

## Q-Learning Agent Report

Musa Mohammadi

H00300499

After running my tests 50 times, I observed the following:

- My tests in total take around 18-20 seconds to converge, this is because the time taken to formulate an optimal policy can vary.
- It requires a minimum of around 30,000 episodes to formulate an effective optimal policy.
- My QI agent wins most of the time, but it never loses. The agent it draws against most of the time is the defensive agent.
- It never loses against a human player; it only wins or draws.